

CE4046 Intelligent Agent Assignment 1

Tianrun Hu

February 22, 2024

1 Method of value iteration

1.1 Description of Implemented Solution

Initialization: The environment is passed into the function by a class that contains all the walls and rewards. A SummaryWriter object from tensorboardX is created for logging purposes. A utility values array V is initialized with zeros for each state in the environment.

Iteration: The update step in the value iteration algorithm is:

$$V_{k+1}(s) = \max_a \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V_k(s')]$$

where $V_k(s)$ is the utility of state s at iteration k , \max_a denotes the maximum value over all possible actions a , $P(s'|s, a)$ is the probability of transitioning to state s' from state s to state s' via action a , γ is the discount factor which prioritizes immediate rewards over distant rewards, $V_k(s')$ is the utility of state s' at iteration k .

Keep tracking the maximum change in utility values δ until it falls below a specified threshold, indicating convergence.

Termination: The iteration loop continues until δ is less than the threshold, at which point it is assumed the utility values have converged to their optimal values.

Plotting: Plot the values and policies calculated from the previous algorithm.

1.2 Plot of Optimal Policy

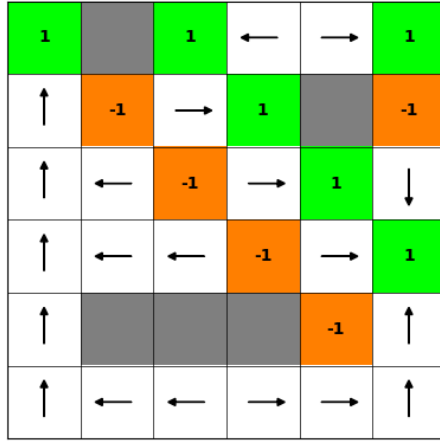


Figure 1: Optimal policy for the gridworld environment. The arrows indicate the action to be taken in each state. The color of the arrows represent the value of the action.

1.3 Utility of all states

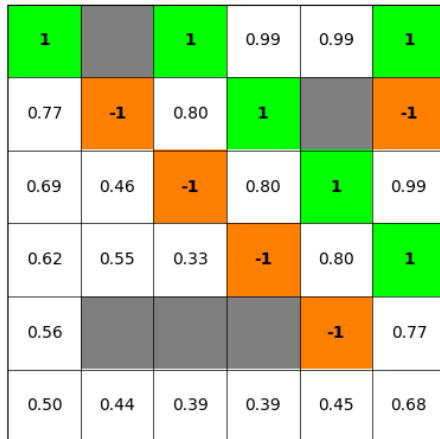


Figure 2: Utility of all states in the gridworld environment. The color of the cells represent the value of the state.

The more detailed value for the final utility is:

0.0	0.0	0.0	0.99445061	0.98753071	0.0
0.77247503	0.0	0.8	0.0	0.0	0.0
0.68529708	0.4611854	0.0	0.8	0.0	0.99445061
0.61840233	0.54986211	0.33239009	0.0	0.8	0.0
0.56078941	0.0	0.0	0.0	0.0	0.77247503
0.49686767	0.4406518	0.38504299	0.3947778	0.45040336	0.68411722

1.4 Convergence of value iteration

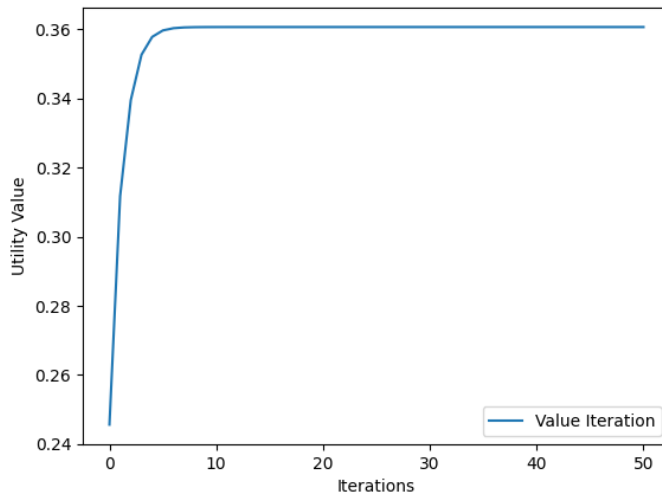


Figure 3: Convergence of value iteration. The y-axis represents the maximum change in utility values across all states. The x-axis represents the number of iterations.

2 Method of policy iteration

2.1 Description of Implemented Solution

Initialization: The setting is basically same with value iteration. Except here we initialize the policy as all actions are set to move right (0, 1).

Policy Evaluation: For each state, the expected utility is computed based on the current policy's action, considering the probability of transitions to all possible next states and the corresponding rewards.

Policy iteration: The iteration consists of two steps: evaluation and improvement.

Policy Evaluation:

$$V^\pi(s) = \sum_{s', r} P(s', r | s, \pi(s)) [r + \gamma V^\pi(s')]$$

where $V^\pi(s)$ is the utility of state s under policy π , $P(s', r|s, \pi(s))$ is the probability of transitioning to state s' with reward r from state s when following policy π . γ is the discount factor and $\pi(s)$ is the action prescribed by policy π in state s .

Policy Improvement:

$$\pi'(s) = \underset{a}{argmax} \sum_{s', r} P(s', r | s, a) [r + \gamma V^\pi(s')]$$

where $pi'(s)$ is the updated action for state s under the improved policy π' , $argmax_a$ denotes the action that maximizes the expected utility.

Termination: The algorithm terminates when the policy is stable (The maximum change in utility values across all states δ is below a threshold, indicating that the utilities have converged).

2.2 Plot of Optimal Policy

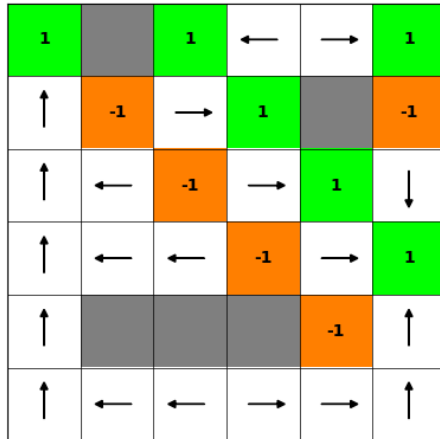


Figure 4: Optimal policy for the gridworld environment. The arrows indicate the action to be taken in each state. The color of the arrows represent the value of the action.

2.3 Utility of all states

1		1	0.99	0.99	1
0.77	-1	0.80	1		-1
0.69	0.46	-1	0.80	1	0.99
0.62	0.55	0.33	-1	0.80	1
0.56				-1	0.77
0.50	0.44	0.39	0.39	0.45	0.68

Figure 5: Utility of all states in the gridworld environment. The color of the cells represent the value of the state.

2.4 Convergence of policy iteration

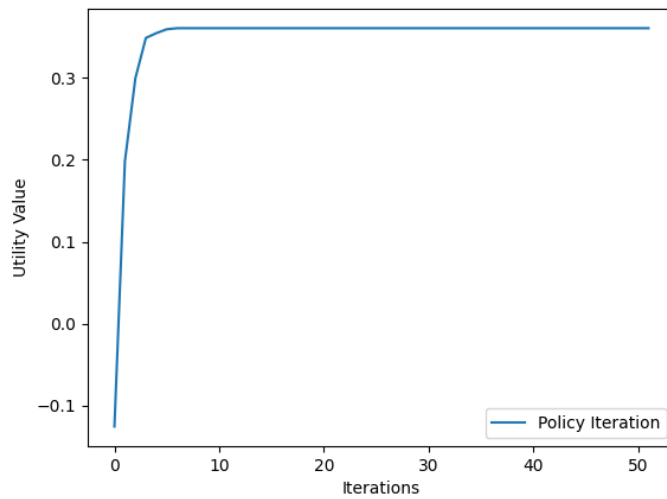


Figure 6: Convergence of policy iteration. The y-axis represents the maximum change in utility values across all states. The x-axis represents the number of iterations.

3 Bonus questions